


# Time Series Learning

# Overview

- Previously: we looked at learning algorithms that did not model time explicitly
- Today: we will look at algorithms that consider time explicitly
  - Filtering & smoothing
  - Time series modeling (ARIMA)
  - Recurrent neural networks




# Time Series

- Time series focus on:
  - Understanding periodicity and trends (including temporal pattern changes)
  - Forecasting
- Time series can be decomposed in three components:
  - Periodic variations (daily, weekly, ... seasonality)
  - Trend (how the mean evolves over time)
  - Irregular variations (left after we remove the periodic variations and trend)

# Stationarity


- A stationary series is one in which the properties – mean, variance and covariance, do not vary with time.




⇒ stationarity

# Stationarity


- A stationary series is one in which the properties – mean, variance and covariance, do not vary with time



mean



Variance



Covariance


⇒ non stationarity

# Stationarity

- Important concept: stationarity
  - If trends and periodic variations are removed
  - Then **variance** of the remaining **residuals** is **constant**
- Prerequisite or intermediate step for many algorithms
- Testing criterion: the lagged ( $\lambda$ ) **autocorrelation** should remain constant



$$r_{\lambda} = \frac{\sum_{t=1}^{N-\lambda} (x_t - \bar{x})(x_{t+\lambda} - \bar{x})}{\sum_{t=1}^N (x_t - \bar{x})^2}$$

Note:  $x_t$  represents the value of one attribute





# Stationarity



- Autocorrelation represents in how far there is a correlation between a time series and a shifted version of itself (with  $\lambda$  steps) => **low autocorrelation**
- **Stationary** if the lagged **autocorrelation** remain constant



random



some memory over  
past time points




cumulative sum of (a)  
– **non-stationary**

# Filtering & smoothing

- Let us assume our time series of values  $x_t$  with a fixed step size  $\Delta t$
- We can apply a filter to our data, taking  $q$  points in the future and past into account:

$$z_t = \sum_{r=-q}^q a_r x_{t+r}$$



- This generates a new time series  $z_t$



# Filtering & smoothing

- What could  $a_r$  look like?
  - If  $a_r = (2q + 1)^{-1}$  it is just the *moving average*
  - If measurements closer to  $t$  are more important we can use *a triangular shape*:

$$a_r = \begin{cases} \frac{q - |r|}{q^2} & -q \leq r \leq q \\ 0 & \text{otherwise} \end{cases}$$

- Or *exponential smoothing*:  $z_t = \alpha^{\text{original}} x_t + (1 - \alpha) z_{t-1}$   
(only past time points mostly)


$$a_r = \alpha(1 - \alpha)^{|r|}$$

$\alpha \sim 1$  : no smoothing

$\alpha \sim 0$  : heavy smoothing

# Filtering & smoothing

- Example exponential smoothing



# Filtering & smoothing: differencing

- Now how can we remove a trend?
- Let us take a filter again, but a simple one:


$$z_t = x_t - x_{t-1} = \nabla x_t$$

- This takes the difference between the current and previous measurement
  - A long term trend has more or less the same influence on the previous and current time point
- We can apply this operator  $d$  times (e.g.  $d=2$ ):  $\nabla^2 x_t = \nabla x_t - \nabla x_{t-1}$

# Filtering & smoothing

- But  $x_{t-1}$  might not be a good estimation of the trend
- We can alternatively use an exponential smoothing  $z_t$  and take  $x_t - z_t$


# Filtering & smoothing




**Fig. 8.4** Black solid line = example time series, red dashed = trends through exponential smoothing, blue dotted line = detrended time series

Detrending w/ exponential smoothing:  $x_t - z_t$  where  $z_t = \alpha * x_t + (1-\alpha) * z_{t-1}$


# A time series can be decomposed



# Time series prediction?



Forecasting my next heart rate?



Forecasting stock prices?

# ARIMA

- **ARIMA**: Auto Regressive Integrated Moving Average
  - A class of models that ‘explains’ a given time series based on
    - AR* → its own past values, that is, its own lags
    - MA* → the lagged forecast errors,
    - so that equation can be used to forecast future values



# ARIMA

- **Auto Regressive (AR only) model**
  - Current value is explained by its **own past values**, that is, its own lags

$$Y_t = \alpha + \beta_1 Y_{t-1} + \beta_2 Y_{t-2} + \dots + \beta_p Y_{t-p} + \epsilon_t$$

- **Moving Average (MA only) model**
  - Current value is explained by the **lagged forecast errors**

$$Y_t = \alpha + \epsilon_t + \phi_1 \epsilon_{t-1} + \phi_2 \epsilon_{t-2} + \dots + \phi_q \epsilon_{t-q}$$

# ARIMA

## ARIMA forecasting equation

- Let  $Y$  denote the *original* series
- Let  $y$  denote the *differenced* (stationarized) series

No difference  $(d=0)$ :  $y_t = Y_t$

First difference  $(d=1)$ :  $y_t = Y_t - Y_{t-1}$

Second difference  $(d=2)$ :  $y_t = (Y_t - Y_{t-1}) - (Y_{t-1} - Y_{t-2})$   
 $= Y_t - 2Y_{t-1} + Y_{t-2}$

Note that the second difference is not just the change relative to two periods ago, i.e., it is *not*  $Y_t - Y_{t-2}$ . Rather, it is the change-in-the-change, which is a measure of local “acceleration” rather than trend.

# ARIMA

## Forecasting equation for $y$

$$\hat{y}_t = \underbrace{\mu}_{\text{constant}} + \underbrace{\phi_1 y_{t-1} + \dots + \phi_p y_{t-p}}_{\text{AR terms (lagged values of } y)} - \underbrace{\theta_1 e_{t-1} \dots - \theta_q e_{t-q}}_{\text{MA terms (lagged errors)}}$$

By convention, the AR terms are + and the MA terms are -

Not as bad as it looks! Usually  $p+q \leq 2$  and either  $p=0$  or  $q=0$  (pure AR or pure MA model)

# ARIMA

## Undifferencing the forecast

The differencing (if any) must be *reversed* to obtain a forecast for the original series:


$$\text{If } d = 0: \quad \hat{Y}_t = \hat{y}_t$$

$$\text{If } d = 1: \quad \hat{Y}_t = \hat{y}_t + Y_{t-1}$$

$$\text{If } d = 2: \quad \hat{Y}_t = \hat{y}_t + 2Y_{t-1} - Y_{t-2}$$

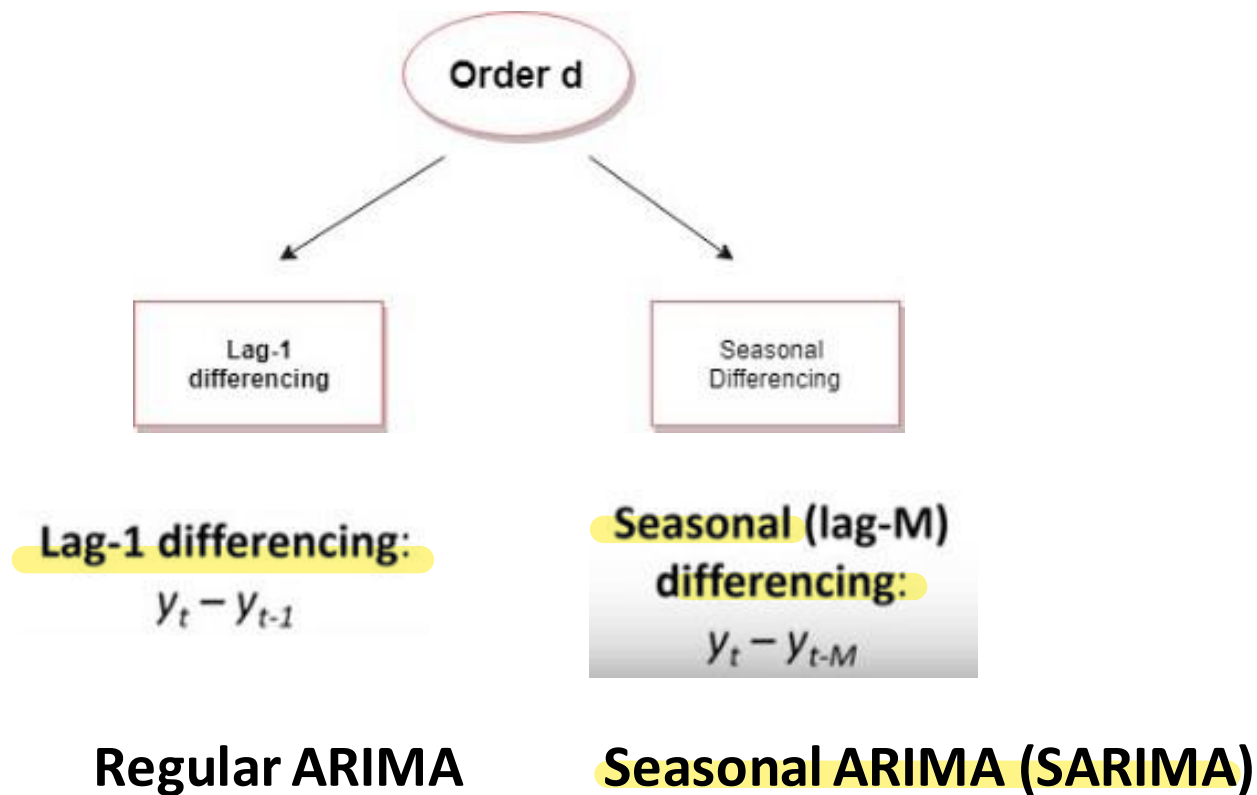
# ARIMA

- Seasonal data requires “seasonal differencing”:  $Y_t - Y_{t-m}$



# ARIMA

- Trend vs. seasonal effects → seasonal time series needs a special care: Seasonal ARIMA




# ARIMA - example

- Let's analyze a fragment of 4000 accelerometer data measurements ( $\approx 20$  s) which we evenly space at a 10ms level to accommodate for a proper time series application



# ARIMA – p & q




**Fig. 8.7** Autocorrelation Function (ACF) for a set of 2000 raw data measurements ( $\approx 10$  s) and Partial Autocorrelation Function (PACF) for the same set of raw data



# ARIMA - example (1)


- One step ahead prediction ( $p=3$ ,  $q=2$ )



The *blue line* represents 500 measurements of the original data, the *red line* the one-step ahead-prediction

# ARIMA - example (2)


- Multiple steps ahead prediction ( $p=3$ ,  $q=2$ )



The blue line indicates the original time series that was used to estimate the ARMA model.<sup>28</sup>  
The red line is the “long-term” prediction, the shaded area its uncertainty


# ARIMA - example (3)

- Seasonality decomposition

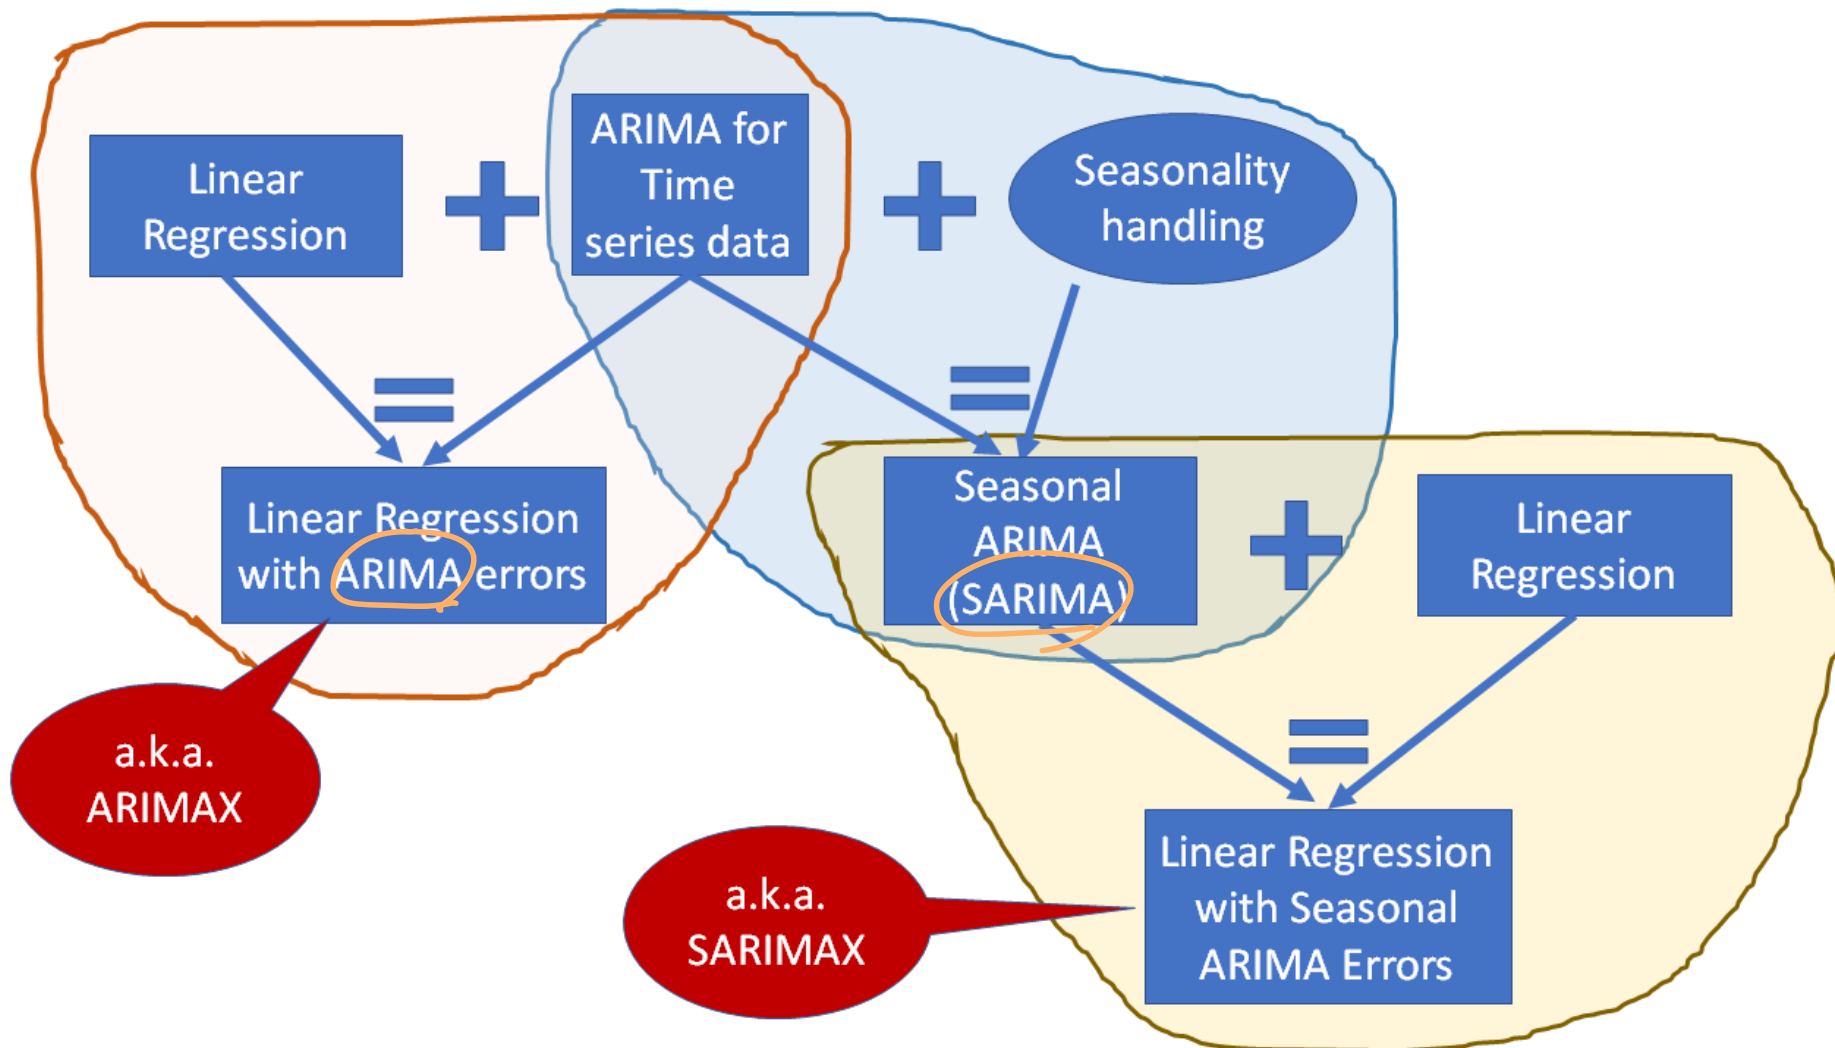


# ARIMA - example (4)

- Multiple steps ahead prediction with seasonality ( $p=3$ ,  $q=2$ ) w/ **S-ARIMA**




# ARIMAX: with Explanatory Variables



# Using a neural network for time-series prediction?

- But a neural network does not consider temporal characteristics of the data




A simple neural net that inputs three items and outputs the next item

# Recurrent neural network



## Recurrent Neuron Network


It looks very much like a feedforward neural network, except it also has connections pointing backward (very simple one with only one input & output w/ recurring output as an input)




## Unrolling the RNN through time

At each *time step*  $t$  (also called a *frame*), this *recurrent neuron* receives the inputs  $x_{(t)}$  as well as its own output from the previous time step,  $y_{(t-1)}$

# Recurrent neural network



RNN with a layer of recurrent neurons



Unrolling the RNN through time


- You can easily create a **layer of recurrent neurons** (= # hidden units)
- At each time step  $t$ , every neuron receives both the input vector  $x(t)$  and the output vector from the previous time step  $y(t-1)$
- Note that both the inputs and outputs are vectors now (when there was just a single neuron, the output was a scalar)



# RNN in Keras




```
model = keras.models.Sequential([  
    keras.layers.SimpleRNN(1, input_shape=[None, 1])  
])
```




```
model = keras.models.Sequential([  
    keras.layers.SimpleRNN(5, input_shape=[None, 1])  
])
```

# Recurrent neural network



Deep RNN: stacking  
three layers of simple  
recurrent neurons




Unrolling the RNN through time

```
model = keras.models.Sequential([  
    keras.layers.SimpleRNN(1, return_sequences=True, input_shape=[None, 1]),  
    keras.layers.SimpleRNN(1, return_sequences=True),  
    keras.layers.SimpleRNN(1)  
])
```


# Case study

- CrowdSignals: Heart rate prediction?



# Case study

- Recall from last time: how did we tune the parameters?
  - Cross validation
  - Does that make sense now?
  - Nope it does not.....




# ARIMAX vs. RNN

- What techniques do we use?

Algorithm	Variant description	Parameters varied
Recurrent Neural Network (RNN)	Recurrent neural network with one layer of hidden neurons with a sigmoid activation function and sigmoid output nodes	number of hidden neurons: {50, 100} maximum iterations over the entire dataset: {250, 500}
Time series	ARIMAX algorithm using Bayesian inference	p: {0, 1, 3, 5} q: {0, 1, 3, 5} d: {0, 1}


# Results

- Actual versus predicted values for ARIMAX



# Results

- Actual versus predicted values for RNN *→ better*



# Summary

- Time series data handling - must consider stationarity (possible to check it with autocorrelation)
- Filtering & smoothing
  - Moving averaging, exponential averaging
  - De-trending with differencing (previous steps or exponential averages)
- ARIMA (auto-regression + moving average of errors)
  - S-ARIMA (seasonal effects) + ARIMAX (explanatory variables: features)
- Recurrent neural networks consider temporal aspects
  - Diverse layer configurations are feasible (e.g., # hidden units, layers, dropouts, statefulness)