



Church-Turing Thesis



The Church-Turing Thesis

Alonzo Church (1903 – 1995) was an American mathematician and logician who made major contributions to mathematical logic and the foundations of theoretical computer science. Turing was his student at Princeton.



- How powerful are TMs?
- What problems can you solve with a computer?
- What does it mean to solve a problem?
 - Rethinking what “solving” a problem means, and two possible answers to that question

Real and “Ideal” Computers

- A real computer has memory limitations: you have a finite amount of RAM, a finite amount of disk space, etc
- However, as computers get more and more powerful, the amount of memory available keeps increasing
- An idealized computer is like a regular computer, but with unlimited RAM and disk space
- It functions just like a regular computer, but never runs out of memory.

Theorem

Turing machines are equal in power to idealized computers. That is, any computation that can be done on a TM can be done on an idealized computer and vice-versa.

Key Idea: Two models of computation are equally powerful if they can simulate each other.

Simulating a TM

- The individual commands in a TM are simple and perform only basic operations:
 - Move*
 - Write*
 - Goto*
 - Return*
 - If*
- The memory for a TM can be thought of as a string with some number keeping track of the current index
- To simulate a TM, we need to
 - see which line of the program we're on
 - determine what command it is
 - simulate that single command
- Claim: This is reasonably straightforward to do on an idealized computer

Simulating a TM

- Because a computer can simulate each individual TM instruction, a computer can do anything a TM can do
- Key Idea: Even the most complicated TM is made out of individual instructions, and if we can simulate those instructions, we can simulate an arbitrarily complicated TM



Simulating Computer



Simulating a Computer

- Programming languages provide a set of simple constructs
 - Think things like variables, arrays, loops, functions, classes, etc.
- You, the programmer, then combine these basic constructs together to assemble larger programs
- Key Idea: If a TM is powerful enough to simulate each of these individual pieces, it's powerful enough to simulate anything a real computer can do

Can TMs Do: Loops?

- We've seen TMs use loops to solve problems
 - The TM for $L = \{ a^n b^n \mid n \in \mathbb{N} \}$ repeatedly pulls off the first and last character from the string
 - Our sorting TM repeatedly finds ba and replaces it with ab
- In some sense, the existence of Goto and labels means that TMs have loops
- Hopefully, it's not too much of a stretch to think that TMs can do while loops, for loops, etc.

Can TMs Do: Arithmetic?

- TMs can perform basic arithmetic
 - Addition of two numbers
 - We can check if two numbers are equal
- We could also do addition and subtraction, compute powers of numbers, do ceilings and floors, etc

Can TMs Do: Variables?

- TMs can maintain variables

- You can think of the TM for $L = \{ a^n b^n \mid n \in \mathbb{N} \}$ as storing two variables - one that counts a number of a's, and one that counts a number of b's
- The TM for Fibonacci numbers tracks the last two Fibonacci numbers, plus the length of the input string

Can TMs Do: Helper Functions?

- We've seen TMs with helper functions
 - We saw how to check for equal numbers of a's and b's by first sorting the string, then checking if the string has the form $a^n b^n$
 - We can check if a decimal number is a Fibonacci number by converting it to unary, then running our unary Fibonacci checker
- A TM could have multiple “helper functions” that work together to solve some larger problem.

What Else Can TMs Do?

- Maintain strings and arrays
 - Store their elements separated with some special separator character
- Support pointers
 - Maintain an array of what's in memory, where each item is tagged with its "memory address."
- Support function call and return
 - It's hard, but you can do this if you can do helper functions and variables

A TM Can Do What Computers Do

- Internally, computers execute by using basic operations like
 - Simple arithmetic
 - Memory reads and writes
 - Branches and jumps
 - Register operations
 - Etc.
- Each of these are simple enough that they could be simulated by a Turing machine.

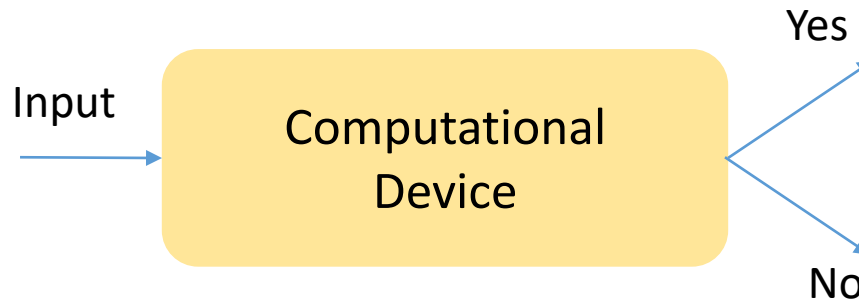


TM \equiv Computer



TM \equiv Computer

- Claim: A TM is powerful enough to simulate any computer program that gets an input, processes that input, then returns some result



- The resulting TM might be colossal, or really slow, or both, but it would still faithfully simulate the computer
- We're going to take this as an article of faith

TM Can Work with

- Images

- A picture is just a 2D array of colors, and a color can be represented as a series of numbers

- Video

- Just a series of pictures

- Music

- Easy

- AI

- Symbol manipulation

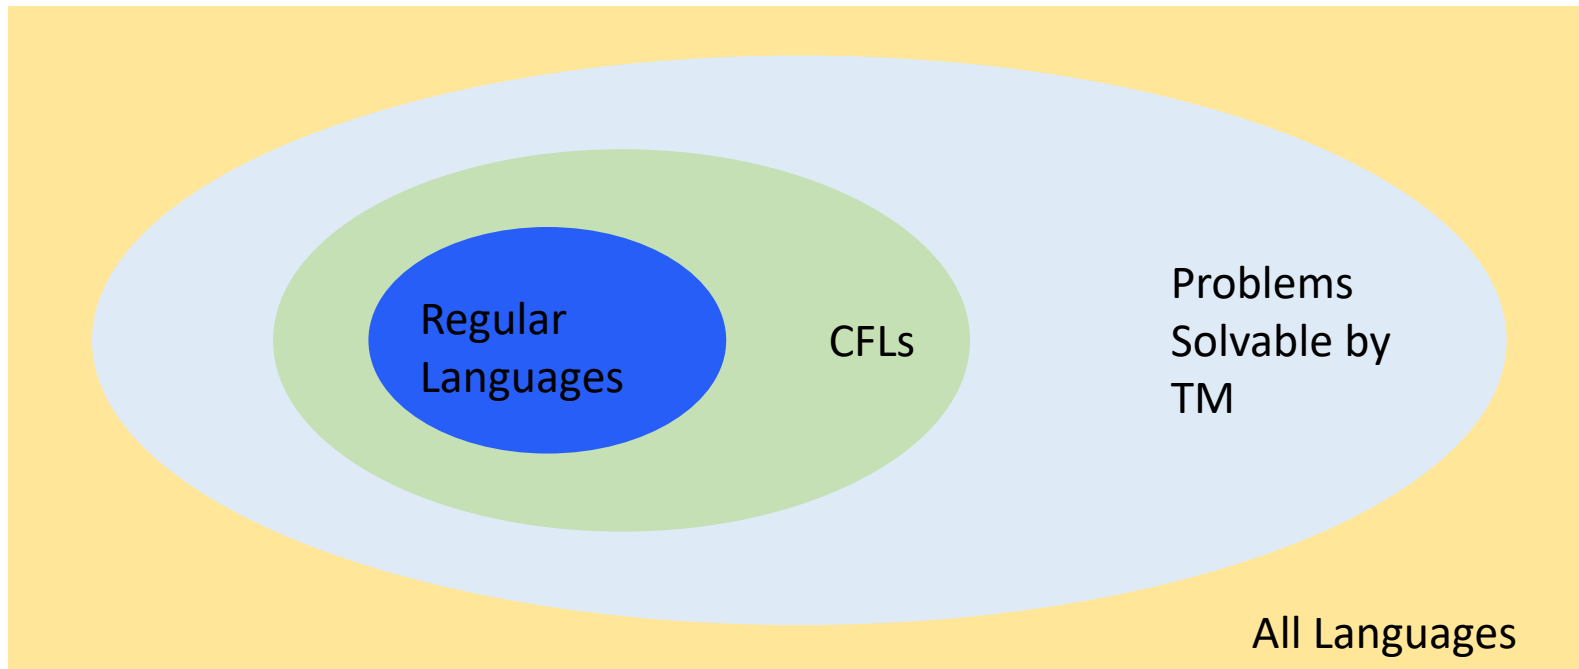
- Deep Learning

- That's just applying a bunch of matrices and nonlinear functions to some input

- An *effective method of computation* is a form of computation with the following properties:
 - The computation consists of a set of steps
 - There are fixed rules governing how one step leads to the next
 - Any computation that yields an answer does so in finitely many steps
 - Any computation that yields an answer always yields the correct answer
- This is not a formal definition. Rather, it's a set of properties we expect out of a computational system

Church-Turing Thesis

- Every effective method of computation is either equivalent to or weaker than a Turing machine.
- “This is not a theorem – it is a falsifiable scientific hypothesis. And it has been thoroughly tested!”



TMs and Computation

- Because Turing machines have the same computational powers as regular computers, we can (essentially) reason about Turing machines by reasoning about actual computer programs
- Going forward, we're going to switch back and forth between TMs and computer programs based on whatever is most appropriate
- In fact, our eventual proofs about the existence of impossible problems will involve a good amount of pseudocode



Finite/Infinite Loops



What problems can we *solve* with a *computer* ?

What does it mean to
“solve” a problem?


What kind of computer?

The Hailstone Sequence

- Consider the following procedure, starting with some $n \in \mathbb{N}$, where $n > 0$:
 - If $n = 1$, you are done
 - If n is even, set $n = n / 2$
 - Otherwise, set $n = 3n + 1$
 - Repeat
- Question: Given a natural number $n > 0$, does this process terminate?


The Hailstone Sequence

● Example: $n = 11$



● Examples:

- $n = 5$? Yes, after 5 steps
- $n = 20$? Yes, after 7 steps
- $n = 7$? Yes, after 16 steps
- $n = 27$? (after 111 steps)



The Hailstone Turing Machine

- Let $\Sigma = \{a\}$ and consider the language $L = \{a^n \mid n > 0 \text{ and the hailstone sequence terminates for } n\}$

Could we build a TM for L ?

- We can build a TM that works as follows:
 - If the input is ϵ , reject
 - While the string is not **a**:
 - If the input has even length, halve the length of the string
 - If the input has odd length, triple the length of the string and append one **a**
 - Accept

The Collatz Conjecture

- It is unknown whether this process will terminate for all natural numbers
- In other words, no one knows whether the TM described before will always stop running!
- The conjecture (unproven claim) that the hailstone sequence always terminates is called the *Collatz Conjecture*
- This problem has eluded a solution for a long time
- Paul Erdős said “Mathematics may not be ready for such problems.”

Paul Erdős (1913 – 1996) was a renowned Hungarian mathematician. He was one of the most prolific mathematicians and producers of mathematical conjectures of the 20th century.

Terence Tao [FAA](#) [FRS](#) (1975) is an Australian-American mathematician. He received the Fields Medal (2006) and Breakthrough Prize at Math. (2014). In 2019, Tao proved that almost all Collatz orbits have finite stopping time.



An Important Observation

- Unlike finite automata, which automatically halt after all the input is read, TMs keep running until they explicitly return true or return false
- As a result, it's possible for a TM to run forever without accepting or rejecting
- This leads to several important questions:
 - How do we formally define what it means to build a TM for a language?
 - What implications does this have about problem solving?

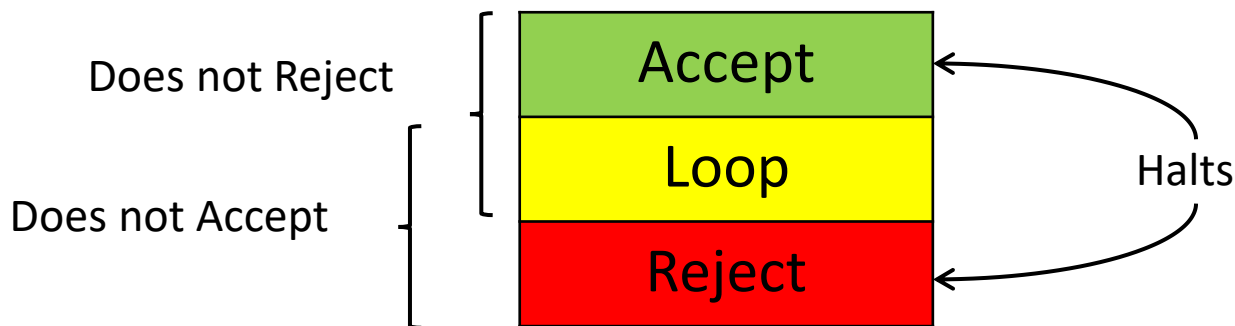


Recognizers and Recognizability



Terminology

- Let M be a Turing machine
- M **accepts** a string w if it returns true on w
- M **rejects** a string w if it returns false on w
- M **loops infinitely** (or **loops**) on a string w if when run on w it neither returns true nor returns false
- M **does not accept w** if it either rejects w or loops on w
- M **does not reject w** if it either accepts w or loops on w
- M **halts on w** if it accepts w or rejects w



Recognizers and Recognizability

- A TM M is called a recognizer for a language L over Σ if the following statement is true:

$$\forall w \in \Sigma^*. (w \in L \leftrightarrow M \text{ accepts } w)$$

- If you are absolutely certain that $w \in L$, then running a recognizer for L on w will (eventually) confirm this
 - Eventually, M will accept w
- If you don't know whether $w \in L$, running M on w may never tell you anything
 - M might loop on w
 - but you can't differentiate between "it'll never give an answer" and "just wait a bit more."
- Does that feel like "solving a problem" to you?

Recognizers and Recognizability

- The hailstone TM M we saw earlier is a recognizer for the language

$$L = \{ a^n \mid n > 0 \text{ and the hailstone sequence terminates for } n \}$$
- If the sequence does terminate starting at n , then M accepts a^n
- If the sequence doesn't terminate, then M loops forever on a^n and never gives an answer
- If you somehow knew the hailstone sequence terminated for n , this machine would (eventually) confirm this
- If you didn't know, this machine might not tell you anything.

Recognizer: Examples

```
bool negatives (string input) {
    return false;
}
```

```
bool positives (string input) {
    return true;
}
```

```
bool infinites (string input) {
    while (true) {
        // do nothing
    }
    return false;
}
```

```
bool repeaters (string input) {
    if (input.size() % 2 != 0) return false;
    for (int i = 0; i < input.size() / 2; i++) {
        if (input[2 * i] != input[2 * i + 1]) {
            return false;
        }
    }
    return true;
}
```

- Each of these code is a recognizer for some language. What language does each recognizer recognize?

Recognizers and Recognizability

- Sums of three cubes

- Are there integers x , y , and z where...

$$x^3 + y^3 + z^3 = 10? \quad \text{Yes! } x = 2, y = 1, z = 1$$

$$x^3 + y^3 + z^3 = 11? \quad \text{Yes! } x = 3, y = -2, z = -2$$

$$x^3 + y^3 + z^3 = 12? \quad \text{Yes! } x = 7, y = 10, z = -11$$

$$x^3 + y^3 + z^3 = 13?$$

Recognizers and Recognizability

- Surprising fact: until 2019, no one knew whether there were integers x , y , and z where $x^3 + y^3 + z^3 = 33$
- A heavily optimized computer search found this answer:
 $x = 8,866,128,975,287,528$
 $y = -8,778,405,442,862,239$
 $z = -2,736,111,468,807,040$
- As of November 2021, no one knows whether there are integers x , y , and z where $x^3 + y^3 + z^3 = 114$.

Recognizers and Recognizability

- Consider the language $L = \{ a^n \mid \exists x \in \mathbb{Z}. \exists y \in \mathbb{Z}. \exists z \in \mathbb{Z}. x^3 + y^3 + z^3 = n \}$
- Here's pseudocode for a recognizer to see whether such a triple exists:


```

for max = 0, 1, 2, ...
  for x from -max to +max:
    for y from -max to +max:
      for z from -max to +max:
        if  $x^3 + y^3 + z^3 = n$ : return true
      
```
- If you somehow knew there was a triple x , y , and z where $x^3 + y^3 + z^3 = n$, running this program will (eventually) convince you of this
- If you weren't sure whether a triple exists, this recognizer might not be useful to you

Recognizers and Recognizability

- The class RE consists of all recognizable languages
- Formally speaking:

$$RE = \{ L \mid L \text{ is a language and there's a recognizer for } L \}$$

- You can think of RE as “all problems with yes/no answers where “yes” answers can be confirmed by a computer.”
 - Given a recognizable language L and a string $w \in L$, running a recognizer for L on w will eventually confirm $w \in L$
 - The recognizer will never have a “false positive” of saying that a string is in L when it isn't
- This is a “weak” notion of solving a problem
- What is a “stronger” one?




Deciders and Decidability



Deciders and Decidability

- Some, but not all, TMs have the following property: the TM halts on all inputs
- If you are given a TM M that always halts, then for the TM M , the statement “ M does not accept w ” means “ M rejects w .”



Deciders and Decidability

- A TM M is called a **decider** for a language L over Σ if the following statements are true:

$$\forall w \in \Sigma^*. M \text{ halts on } w$$

$$\forall w \in \Sigma^*. (w \in L \leftrightarrow M \text{ accepts } w)$$

- In other words, M accepts all strings in L and rejects all strings not in L
- In other words, M is a recognizer for L , and M halts on all inputs
- If you aren't sure whether $w \in L$, running M on w will (eventually) give you an answer to that question.

Deciders and Decidability

- The hailstone TM M we saw earlier is a recognizer for the language

$$L = \{ a^n \mid n > 0 \text{ and the hailstone sequence terminates for } n \}$$

- If the hailstone sequence terminates for n , then M accepts a^n . If it doesn't, then M does not accept a^n
- We honestly don't know if M is a decider for this language
 - If the hailstone sequence always terminates, then M always halts and is a decider for L
 - If the hailstone sequence doesn't always terminate, then M will loop on some inputs and isn't a decider for L

Recognizer: Examples

```
bool negatives (string input) {
    return false;
}
```

```
bool positives (string input) {
    return true;
}
```

```
bool infinites (string input) {
    while (true) {
        // do nothing
    }
    return false;
}
```

```
bool repeaters (string input) {
    if (input.size() % 2 != 0) return false;
    for (int i = 0; i < input.size() / 2; i++) {
        if (input[2 * i] != input[2 * i + 1]) {
            return false;
        }
    }
    return true;
}
```

- Each of these code is a recognizer for some language. Which are deciders?

Deciders and Decidability

- While no one knows whether there are integers x , y , and z where $x^3 + y^3 + z^3 = 114$, it is very easy to figure out whether there are integers x , y , and z where $x^2 + y^2 + z^2 = 114$
- Why?

Deciders and Decidability

- Consider the language $L = \{a^n \mid \exists x \in \mathbb{Z}. \exists y \in \mathbb{Z}. \exists z \in \mathbb{Z}. x^2 + y^2 + z^2 = n\}$
- Here's pseudocode for a decider to see whether such a triple exists:

```

for x from 0 to n:
  for y from 0 to n:
    for z from 0 to n:
      if  $x^2 + y^2 + z^2 = n$ : return true
return false

```

- After trying all possible options, this program will either find a triple that works or report that none exists

Deciders and Decidability

- The class R consists of all decidable languages
- Formally speaking:

$$R = \{ L \mid L \text{ is a language and there's a decider for } L \}$$
- You can think of R as “all problems with yes/no answers that can be fully solved by computers.”
 - Given a decidable language, run a decider for L and see what happens
 - Think of this as “knowledge creation” – if you don’t know whether a string is in L , running the decider will, given enough time, tell you
- The class R contains all the regular languages, all the context-free languages, most of algorithms, etc.
- This is a “strong” notion of solving a problem.

R and RE Languages

- Every decider for L is also a recognizer for L
- This means that $R \subseteq RE$
- Hugely important theoretical question:

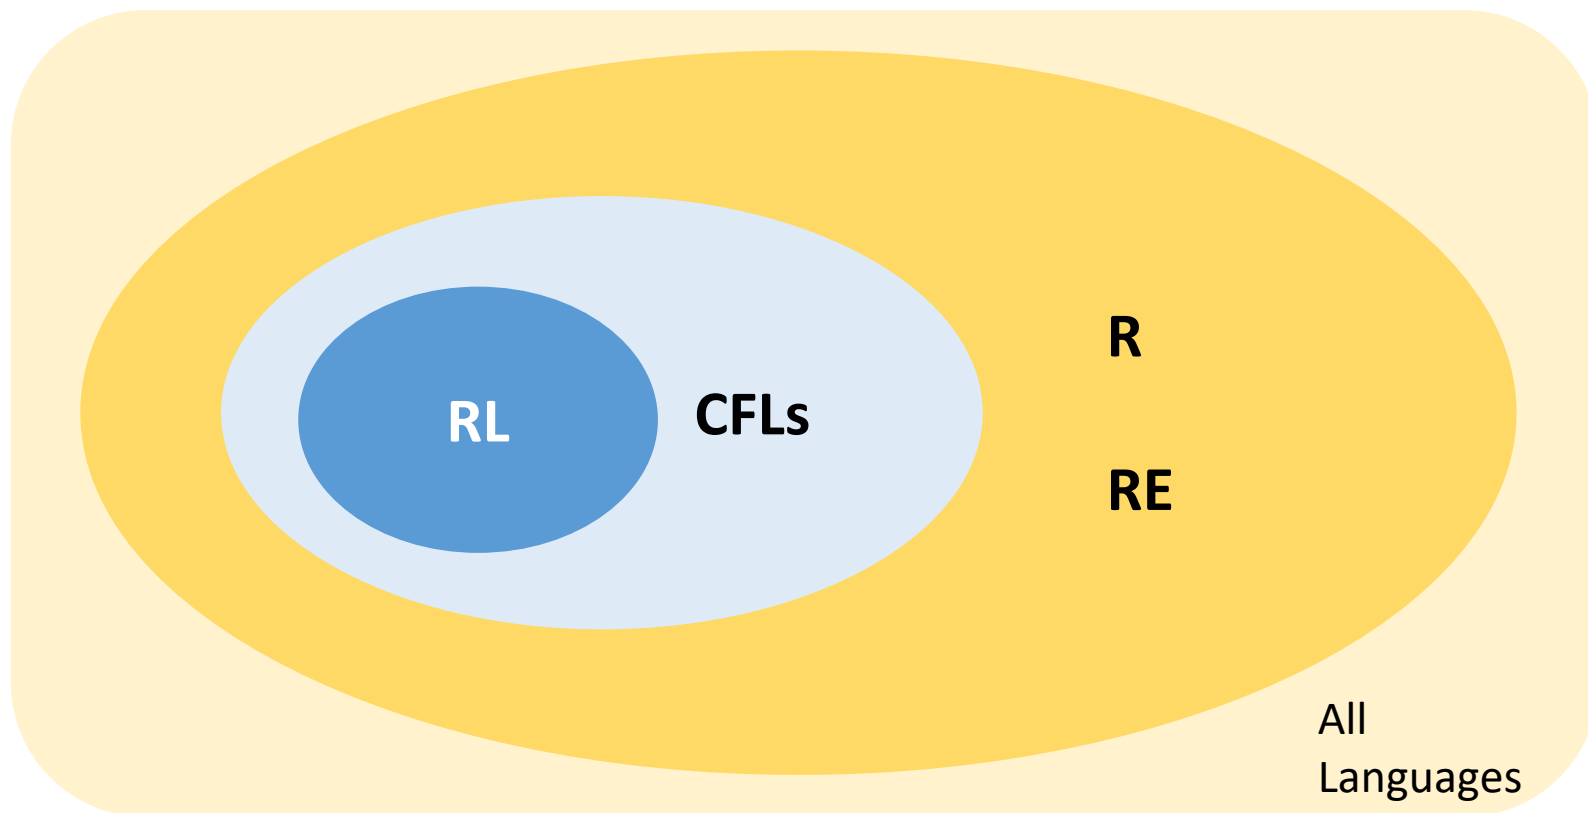
$$R \stackrel{?}{=} RE$$

- That is, if you can just confirm “yes” answers to a problem, can you necessarily solve that problem?

Relations

By Definition, $R \subseteq RE$

$R = RE$???



Questions

- Why exactly is RE an interesting class of problems?
- What does the $R \stackrel{?}{=} RE$ question mean?
- Is $R = RE$?
- What lies beyond R and RE?